

```
In [1]: # importing required libraries
import torch.nn as nn
import torch
import torch.nn.functional as F
import math,copy,re
import warnings
import pandas as pd
import numpy as np
import seaborn as sns
import torchtext
import matplotlib.pyplot as plt
warnings.simplefilter("ignore")
print(torch.__version__)

1.13.1+cu117
```

```
In [2]: class Embedding(nn.Module):
def __init__(self, vocab_size, embed_dim):
    """
    Args:
        vocab_size: size of vocabulary
        embed_dim: dimension of embeddings
    """
    super(Embedding, self).__init__()
    self.embed = nn.Embedding(vocab_size, embed_dim)
def forward(self, x):
    """
    Args:
        x: input vector
    Returns:
        out: embedding vector
    """
    out = self.embed(x)
    return out
```

```
In [3]: # register buffer in Pytorch ->
# If you have parameters in your model, which should be saved and restored in the state_dict,
# but not trained by the optimizer, you should register them as buffers.
```

```
class PositionalEmbedding(nn.Module):
def __init__(self, max_seq_len, embed_model_dim):
    """
    Args:
        seq_len: length of input sequence
        embed_model_dim: dimension of embedding
    """
    super(PositionalEmbedding, self).__init__()
    self.embed_dim = embed_model_dim

    pe = torch.zeros(max_seq_len, self.embed_dim)
    for pos in range(max_seq_len):
        for i in range(0, self.embed_dim, 2):
            pe[pos, i] = math.sin(pos / (10000 ** ((2 * i) / self.embed_dim)))
            pe[pos, i + 1] = math.cos(pos / (10000 ** ((2 * (i + 1)) / self.embed_dim)))
    pe = pe.unsqueeze(0)
    self.register_buffer('pe', pe)

def forward(self, x):
    """
    Args:
        x: input vector
    Returns:
        x: output
    """

    # make embeddings relatively larger
    x = x * math.sqrt(self.embed_dim)
    # add constant to embedding
    seq_len = x.size(1)
    x = x + torch.autograd.Variable(self.pe[:, :seq_len], requires_grad=False)
    return x
```

```
In [4]: class MultiHeadAttention(nn.Module):
def __init__(self, embed_dim=512, n_heads=8):
    """
    Args:
        embed_dim: dimension of embedding vector output
        n_heads: number of self attention heads
    """
    super(MultiHeadAttention, self).__init__()

    self.embed_dim = embed_dim #512 dim
    self.n_heads = n_heads #8
    self.single_head_dim = int(self.embed_dim / self.n_heads) #512/8 = 64 . each key,query, value will be of 64d

    #key,query and value matrixes #64 x 64
    self.query_matrix = nn.Linear(self.single_head_dim, self.single_head_dim, bias=False) # single key matrix for all 8 keys #512x512
    self.key_matrix = nn.Linear(self.single_head_dim, self.single_head_dim, bias=False)
    self.value_matrix = nn.Linear(self.single_head_dim, self.single_head_dim, bias=False)
    self.out = nn.Linear(self.n_heads*self.single_head_dim, self.embed_dim)

def forward(self, key, query, value, mask=None): #batch_size x sequence_length x embedding_dim # 32 x 10 x 512
    """
    Args:
        key : key vector
        query : query vector
        value : value vector
        mask: mask for decoder

    Returns:
        output vector from multihead attention
    """
    batch_size = key.size(0)
    seq_length = key.size(1)

    # query dimension can change in decoder during inference.
    # so we cant take general seq_length
    seq_length_query = query.size(1)

    # 32x10x512
    key = key.view(batch_size, seq_length, self.n_heads, self.single_head_dim) #batch_size x sequence_length x n_heads x single_head_dim = (32x10x8x64)
    query = query.view(batch_size, seq_length_query, self.n_heads, self.single_head_dim) #(32x10x8x64)
    value = value.view(batch_size, seq_length, self.n_heads, self.single_head_dim) #(32x10x8x64)

    k = self.key_matrix(key) # (32x10x8x64)
    q = self.query_matrix(query)
    v = self.value_matrix(value)

    q = q.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim) # (32 x 8 x 10 x 64)
    k = k.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)
    v = v.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)

    # computes attention
    # adjust key for matrix multiplication
    k_adjusred = k.transpose(-1,-2) # (batch_size, n_heads, single_head_dim, seq_len) #(32 x 8 x 64 x 10)
    product = torch.matmul(q, k_adjusred) #(32 x 8 x 10 x 64) x (32 x 8 x 64 x 10) = #(32x8x10x10)

    # fill those positions of product matrix as (-1e20) where mask positions are 0
```

```

if mask is not None:
    product = product.masked_fill(mask == 0, float("-1e20"))

#dividing by square root of key dimension
product = product / math.sqrt(self.single_head_dim) # / sqrt(64)

#applying softmax
scores = F.softmax(product, dim=-1)

#multiply with value matrix
scores = torch.matmul(scores, v)  ##(32x8x 10x 10) x (32 x 8 x 10 x 64) = (32 x 8 x 10 x 64)

#concatenated output
concat = scores.transpose(1,2).contiguous().view(batch_size, seq_length_query, self.single_head_dim*self.n_heads)  # (32x8x10x64) -> (32x10x8x64) -> (32,10,512)

output = self.out(concat)  # (32,10,512) -> (32,10,512)

return output

```

```

In [5]: class TransformerBlock(nn.Module):
def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
    super(TransformerBlock, self).__init__()

    """
    Args:
        embed_dim: dimension of the embedding
        expansion_factor: factor which determines output dimension of linear layer
        n_heads: number of attention heads
    """

    self.attention = MultiHeadAttention(embed_dim, n_heads)

    self.norm1 = nn.LayerNorm(embed_dim)
    self.norm2 = nn.LayerNorm(embed_dim)

    self.feed_forward = nn.Sequential(
        nn.Linear(embed_dim, expansion_factor*embed_dim),
        nn.ReLU(),
        nn.Linear(expansion_factor*embed_dim, embed_dim)
    )

    self.dropout1 = nn.Dropout(0.2)
    self.dropout2 = nn.Dropout(0.2)

def forward(self, key, query, value):
    """
    Args:
        key: key vector
        query: query vector
        value: value vector
        norm2_out: output of transformer block
    """

    attention_out = self.attention(key, query, value)  #32x10x512
    attention_residual_out = attention_out + value  #32x10x512
    norm1_out = self.dropout1(self.norm1(attention_residual_out))  #32x10x512

    feed_fwd_out = self.feed_forward(norm1_out)  #32x10x512 -> #32x10x2048 -> 32x10x512
    feed_fwd_residual_out = feed_fwd_out + norm1_out  #32x10x512
    norm2_out = self.dropout2(self.norm2(feed_fwd_residual_out))  #32x10x512

    return norm2_out

class TransformerEncoder(nn.Module):
    """
    Args:
        seq_len : length of input sequence
        embed_dim: dimension of embedding
        num_layers: number of encoder layers
        expansion_factor: factor which determines number of linear layers in feed forward layer
        n_heads: number of heads in multihead attention

    Returns:
        out: output of the encoder
    """
def __init__(self, seq_len, vocab_size, embed_dim, num_layers=2, expansion_factor=4, n_heads=8):
    super(TransformerEncoder, self).__init__()

    self.embedding_layer = Embedding(vocab_size, embed_dim)
    self.positional_encoder = PositionalEncoding(seq_len, embed_dim)

    self.layers = nn.ModuleList([TransformerBlock(embed_dim, expansion_factor, n_heads) for i in range(num_layers)])

def forward(self, x):
    embed_out = self.embedding_layer(x)
    out = self.positional_encoder(embed_out)
    for layer in self.layers:
        out = layer(out, out, out)

    return out  #32x10x512

```

```

In [6]: class TransformerBlock(nn.Module):
def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
    super(TransformerBlock, self).__init__()

    """
    Args:
        embed_dim: dimension of the embedding
        expansion_factor: factor which determines output dimension of linear layer
        n_heads: number of attention heads
    """

    self.attention = MultiHeadAttention(embed_dim, n_heads)

    self.norm1 = nn.LayerNorm(embed_dim)
    self.norm2 = nn.LayerNorm(embed_dim)

    self.feed_forward = nn.Sequential(
        nn.Linear(embed_dim, expansion_factor*embed_dim),
        nn.ReLU(),
        nn.Linear(expansion_factor*embed_dim, embed_dim)
    )

    self.dropout1 = nn.Dropout(0.2)
    self.dropout2 = nn.Dropout(0.2)

def forward(self, key, query, value):
    """
    Args:
        key: key vector
        query: query vector
        value: value vector
        norm2_out: output of transformer block
    """

    attention_out = self.attention(key, query, value)  #32x10x512
    attention_residual_out = attention_out + value  #32x10x512

```

```

        norm1_out = self.dropout1(self.norm1(attention_residual_out)) #32x10x512

        feed_fwd_out = self.feed_forward(norm1_out) #32x10x512 -> #32x10x2048 -> 32x10x512
        feed_fwd_residual_out = feed_fwd_out + norm1_out #32x10x512
        norm2_out = self.dropout2(self.norm2(feed_fwd_residual_out)) #32x10x512

        return norm2_out

class TransformerEncoder(nn.Module):
    """
    Args:
        seq_len : length of input sequence
        embed_dim: dimension of embedding
        num_layers: number of encoder layers
        expansion_factor: factor which determines number of linear layers in feed forward layer
        n_heads: number of heads in multihead attention

    Returns:
        out: output of the encoder
    """
    def __init__(self, seq_len, vocab_size, embed_dim, num_layers=2, expansion_factor=4, n_heads=8):
        super(TransformerEncoder, self).__init__()

        self.embedding_layer = Embedding(vocab_size, embed_dim)
        self.positional_encoder = PositionalEmbedding(seq_len, embed_dim)

        self.layers = nn.ModuleList([TransformerBlock(embed_dim, expansion_factor, n_heads) for i in range(num_layers)])

    def forward(self, x):
        embed_out = self.embedding_layer(x)
        out = self.positional_encoder(embed_out)
        for layer in self.layers:
            out = layer(out, out, out)

        return out #32x10x512

```

```

In [7]: class DecoderBlock(nn.Module):
    def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
        super(DecoderBlock, self).__init__()

        """
        Args:
            embed_dim: dimension of the embedding
            expansion_factor: factor which determines output dimension of linear layer
            n_heads: number of attention heads

        """
        self.attention = MultiHeadAttention(embed_dim, n_heads=8)
        self.norm = nn.LayerNorm(embed_dim)
        self.dropout = nn.Dropout(0.2)
        self.transformer_block = TransformerBlock(embed_dim, expansion_factor, n_heads)

    def forward(self, key, query, x, mask):
        """
        Args:
            key: key vector
            query: query vector
            value: value vector
            mask: mask to be given for multi head attention
        Returns:
            out: output of transformer block
        """

        #we need to pass mask mask only to fst attention
        attention = self.attention(x, x, x, mask=mask) #32x10x512
        value = self.dropout(self.norm(attention + x))

        out = self.transformer_block(key, query, value)

        return out

class TransformerDecoder(nn.Module):
    def __init__(self, target_vocab_size, embed_dim, seq_len, num_layers=2, expansion_factor=4, n_heads=8):
        super(TransformerDecoder, self).__init__()
        """
        Args:
            target_vocab_size: vocabulary size of target
            embed_dim: dimension of embedding
            seq_len : length of input sequence
            num_layers: number of encoder layers
            expansion_factor: factor which determines number of linear layers in feed forward layer
            n_heads: number of heads in multihead attention

        """
        self.word_embedding = nn.Embedding(target_vocab_size, embed_dim)
        self.position_embedding = PositionalEmbedding(seq_len, embed_dim)

        self.layers = nn.ModuleList(
            [
                DecoderBlock(embed_dim, expansion_factor=4, n_heads=8)
                for _ in range(num_layers)
            ]
        )
        self.fc_out = nn.Linear(embed_dim, target_vocab_size)
        self.dropout = nn.Dropout(0.2)

    def forward(self, x, enc_out, mask):
        """
        Args:
            x: input vector from target
            enc_out : output from encoder layer
            trg_mask: mask for decoder self attention
        Returns:
            out: output vector
        """

        x = self.word_embedding(x) #32x10x512
        x = self.position_embedding(x) #32x10x512
        x = self.dropout(x)

        for layer in self.layers:
            x = layer(enc_out, x, enc_out, mask)

        out = F.softmax(self.fc_out(x))

        return out

```

```

In [8]: class Transformer(nn.Module):
    def __init__(self, embed_dim, src_vocab_size, target_vocab_size, seq_length, num_layers=2, expansion_factor=4, n_heads=8):
        super(Transformer, self).__init__()

        """
        Args:

```

```

        embed_dim: dimension of embedding
        src_vocab_size: vocabulary size of source
        target_vocab_size: vocabulary size of target
        seq_length : length of input sequence
        num_layers: number of encoder layers
        expansion_factor: factor which determines number of linear layers in feed forward layer
        n_heads: number of heads in multihead attention

    """

    self.target_vocab_size = target_vocab_size

    self.encoder = TransformerEncoder(seq_length, src_vocab_size, embed_dim, num_layers=num_layers, expansion_factor=expansion_factor, n_heads=n_heads)
    self.decoder = TransformerDecoder(target_vocab_size, embed_dim, seq_length, num_layers=num_layers, expansion_factor=expansion_factor, n_heads=n_heads)

def make_trg_mask(self, trg):
    """
    Args:
        trg: target sequence
    Returns:
        trg_mask: target mask
    """
    batch_size, trg_len = trg.shape
    # returns the lower triangular part of matrix filled with ones
    trg_mask = torch.tril(torch.ones((trg_len, trg_len))).expand(
        batch_size, 1, trg_len, trg_len
    )
    return trg_mask

def decode(self,src,trg):
    """
    for inference
    Args:
        src: input to encoder
        trg: input to decoder
    out:
        out_labels : returns final prediction of sequence
    """
    trg_mask = self.make_trg_mask(trg)
    enc_out = self.encoder(src)
    out_labels = []
    batch_size,seq_len = src.shape[0],src.shape[1]
    #outputs = torch.zeros(seq_len, batch_size, self.target_vocab_size)
    out = trg
    for i in range(seq_len): #10
        out = self.decoder(out,enc_out,trg_mask) #bs x seq_len x vocab_dim
        # taking the last token
        out = out[:, -1, :]

        out = out.argmax(-1)
        out_labels.append(out.item())
        out = torch.unsqueeze(out,axis=0)

    return out_labels

def forward(self, src, trg):
    """
    Args:
        src: input to encoder
        trg: input to decoder
    out:
        out: final vector which returns probabilities of each target word
    """
    trg_mask = self.make_trg_mask(trg)
    enc_out = self.encoder(src)

    outputs = self.decoder(trg, enc_out, trg_mask)
    return outputs

```

```

In [9]: src_vocab_size = 11
        target_vocab_size = 11
        num_layers = 6
        seq_length= 12

# let 0 be sos token and 1 be eos token
src = torch.tensor([[0, 2, 5, 6, 4, 3, 9, 5, 2, 9, 10, 1],
                    [0, 2, 8, 7, 3, 4, 5, 6, 7, 2, 10, 1]])
target = torch.tensor([[0, 1, 7, 4, 3, 5, 9, 2, 8, 10, 9, 1],
                       [0, 1, 5, 6, 2, 4, 7, 6, 2, 8, 10, 1]])

print(src.shape,target.shape)
model = Transformer(embed_dim=512, src_vocab_size=src_vocab_size,
                    target_vocab_size=target_vocab_size, seq_length=seq_length,
                    num_layers=num_layers, expansion_factor=4, n_heads=8)

model

torch.Size([2, 12]) torch.Size([2, 12])

```

[illegible]

[illegible]

```

(norm2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
(feed_forward): Sequential(
  (0): Linear(in_features=512, out_features=2048, bias=True)
  (1): ReLU()
  (2): Linear(in_features=2048, out_features=512, bias=True)
)
(dropout1): Dropout(p=0.2, inplace=False)
(dropout2): Dropout(p=0.2, inplace=False)
)
)
)
(fc_out): Linear(in_features=512, out_features=11, bias=True)
(dropout): Dropout(p=0.2, inplace=False)
)
)

```

```
In [10]: out = model(src, target)
out.shape
```

```
Out[10]: torch.Size([2, 12, 11])
```

```
In [11]: # inference
model = Transformer(embed_dim=512, src_vocab_size=src_vocab_size,
                    target_vocab_size=target_vocab_size, seq_length=seq_length,
                    num_layers=num_layers, expansion_factor=4, n_heads=8)

src = torch.tensor([[0, 2, 5, 6, 4, 3, 9, 5, 2, 9, 10, 1]])
trg = torch.tensor([[0]])
print(src.shape, trg.shape)
out = model.decode(src, trg)
out

torch.Size([1, 12]) torch.Size([1, 1])
```

```
Out[11]: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
In [12]: pip install torchtext
```

```

Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: torchtext in /home/ubuntu/.local/lib/python3.10/site-packages (0.14.1)
Requirement already satisfied: numpy in /home/ubuntu/.local/lib/python3.10/site-packages (from torchtext) (1.24.2)
Requirement already satisfied: tqdm in /home/ubuntu/.local/lib/python3.10/site-packages (from torchtext) (4.64.1)
Requirement already satisfied: requests in /home/ubuntu/.local/lib/python3.10/site-packages (from torchtext) (2.28.2)
Requirement already satisfied: torch==1.13.1 in /home/ubuntu/.local/lib/python3.10/site-packages (from torchtext) (1.13.1)
Requirement already satisfied: nvidia-cuda-runtime-cu11==11.7.99 in /home/ubuntu/.local/lib/python3.10/site-packages (from torch==1.13.1->torchtext) (11.7.99)
Requirement already satisfied: nvidia-cuda-nvrtc-cu11==11.7.99 in /home/ubuntu/.local/lib/python3.10/site-packages (from torch==1.13.1->torchtext) (11.7.99)
Requirement already satisfied: nvidia-cudnn-cu11==8.5.0.96 in /home/ubuntu/.local/lib/python3.10/site-packages (from torch==1.13.1->torchtext) (8.5.0.96)
Requirement already satisfied: nvidia-cublas-cu11==11.10.3.66 in /home/ubuntu/.local/lib/python3.10/site-packages (from torch==1.13.1->torchtext) (11.10.3.66)
Requirement already satisfied: typing-extensions in /home/ubuntu/.local/lib/python3.10/site-packages (from torch==1.13.1->torchtext) (4.5.0)
Requirement already satisfied: wheel in /usr/lib/python3/dist-packages (from nvidia-cublas-cu11==11.10.3.66->torch==1.13.1->torchtext) (0.37.1)
Requirement already satisfied: setuptools in /usr/lib/python3/dist-packages (from nvidia-cublas-cu11==11.10.3.66->torch==1.13.1->torchtext) (59.6.0)
Requirement already satisfied: idna<4,>=2.5 in /home/ubuntu/.local/lib/python3.10/site-packages (from requests->torchtext) (2.10)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/lib/python3/dist-packages (from requests->torchtext) (1.26.5)
Requirement already satisfied: certifi>=2017.4.17 in /usr/lib/python3/dist-packages (from requests->torchtext) (2020.6.20)
Requirement already satisfied: charset-normalizer<4,>=2 in /home/ubuntu/.local/lib/python3.10/site-packages (from requests->torchtext) (2.1.1)
Note: you may need to restart the kernel to use updated packages.

```

```
In [ ]:
```

```
In [ ]:
```